Data Input

Module 4

Andrew Jaffe

June 15, 2015

Data Input

- We used several pre-installed sample datasets during previous modules (`CO2`, `iris`)
- However, 'reading in' data is the first step of any real project/analysis
- R can read almost any file format, especially via add-on packages
- We are going to focus on simple delimited files first

    - tab delimited (e.g. '.txt')
    - comma separated (e.g. '.csv')
    - Microsoft excel (e.g. '.xlsx')

Data Input

`read.table()`: Reads a file in table format and creates a data frame from it, with cases corresponding to lines and variables to fields in the file.

```
# the four ones I've put at the top are the important inputs
read.table( file, # filename
            header = FALSE, # are there column names?
            sep = "", # what separates columns?
            as.is = !stringsAsFactors, # do you want character strings as factors or characters?
            quote = "\"'",  dec = ".", row.names, col.names,
            na.strings = "NA", nrows = -1,
            skip = 0, check.names = TRUE, fill = !blank.lines.skip,
            strip.white = FALSE, blank.lines.skip = TRUE, comment.char = "#",
            stringsAsFactors = default.stringsAsFactors())

# for example: `read.table("file.txt", header = TRUE, sep="\t", as.is=TRUE)`
```

Data Input

- The filename is the path to your file, in quotes
- The function will look in your "working directory" if no absolute file path is given
- Note that the filename can also be a path to a file on a website (e.g. 'www.someurl.com/table1.txt')

Data Aside

- Everything we do in class will be using real publicly available data - there are few 'toy' example datasets and 'simulated' data
- OpenBaltimore and Data.gov will be sources for the first few days

Data Input

Monuments Dataset: "This data set shows the point location of Baltimore City monuments. However, the completness and currentness of these data are uncertain."

- Navigate to: https://data.baltimorecity.gov/Community/Monuments/cpxf-kxp3
- Export –> Download –> Download As: CSV
- Save it (or move it) to the same folder as your day1.R script
- Within RStudio: Session –> Set Working Directory –> To Source File Location

Data Input

There is a 'wrapper' function for reading CSV files:

```
read.csv
```

```
## function (file, header = TRUE, sep = ",", quote = "\"", dec = ".",
##     fill = TRUE, comment.char = "", ...)
## read.table(file = file, header = header, sep = sep, quote = quote,
##     dec = dec, fill = fill, comment.char = comment.char, ...)
## <bytecode: 0x00000000126a2208>
## <environment: namespace:utils>
```

Note: the ... designates extra/optional arguments that can be passed to `read.table()` if needed

Data Input

- Starting out, you can use RStudio –> Tools –> Import Dataset –> From Text File and select

```
mon = read.csv("../data/Monuments.csv",header=TRUE,as.is=TRUE)
head(mon)
```

```
##                              name zipCode neighborhood councilDistrict
## 1          James Cardinal Gibbons   21201     Downtown              11
## 2              The Battle Monument   21202     Downtown              11
## 3 Negro Heroes of the U.S Monument   21202     Downtown              11
## 4              Star Bangled Banner   21202     Downtown              11
## 5  Flame at the Holocaust Monument   21202     Downtown              11
## 6                   Calvert Statue   21202     Downtown              11
##   policeDistrict                  Location.1
## 1        CENTRAL  408 CHARLES ST\nBaltimore, MD\n
## 2        CENTRAL
## 3        CENTRAL
## 4        CENTRAL 100 HOLLIDAY ST\nBaltimore, MD\n
## 5        CENTRAL     50 MARKET PL\nBaltimore, MD\n
## 6        CENTRAL  100 CALVERT ST\nBaltimore, MD\n
```

Data Input

```r
colnames(mon)
```

```
## [1] "name"           "zipCode"          "neighborhood"     "councilDistrict"
## [5] "policeDistrict" "Location.1"
```

```r
head(mon$zipCode)
```

```
## [1] 21201 21202 21202 21202 21202 21202
```

```r
head(mon$neighborhood)
```

```
## [1] "Downtown" "Downtown" "Downtown" "Downtown" "Downtown" "Downtown"
```

Aside: Working Directory

- R looks for files on your computer relative to the "working" directory
- It's always safer to set the working directory at the beginning of your script. Note that setting the working directory created the necessary code that you can copy into your script.
- Example of help file

```r
## get the working directory
getwd()
# setwd("~/Dropbox/summerR_2015/Lectures")
```

Aside: Working Directory

- Setting the directory can sometimes be finicky

    - Windows: Default directory structure involves single backslashes (""), but R interprets these as"escape" characters. So you must replace the backslash with forward slashed ("/") or two backslashes ("\")
    - Mac/Linux: Default is forward slashes, so you are okay

- Typical linux/DOS directory structure syntax applies

    - ".." goes up one level
    - "./" is the current directory
    - "~" is your home directory

Working Directory

Try some directory navigation:

```r
dir("./") # shows directory contents
```

```
##  [1] "module1.html"   "module1.pdf"    "module1.R"
##  [4] "module1.Rmd"    "module10.html"  "module10.pdf"
##  [7] "module10.R"     "module10.Rmd"   "module10_cache"
## [10] "module10_files" "module11.html"  "module11.pdf"
```

```
## [13] "module11.R"      "module11.Rmd"    "module11_cache"
## [16] "module11_files"  "module2.html"    "module2.pdf"
## [19] "module2.R"       "module2.Rmd"     "module3.html"
## [22] "module3.pdf"     "module3.R"       "module3.Rmd"
## [25] "module4.html"    "module4.pdf"     "module4.R"
## [28] "module4.Rmd"     "module5.html"    "module5.pdf"
## [31] "module5.R"       "module5.Rmd"     "module6.html"
## [34] "module6.pdf"     "module6.R"       "module6.Rmd"
## [37] "module7.html"    "module7.pdf"     "module7.R"
## [40] "module7.Rmd"     "module8.html"    "module8.pdf"
## [43] "module8.R"       "module8.Rmd"     "module9.html"
## [46] "module9.pdf"     "module9.R"       "module9.Rmd"
## [49] "module9_files"   "renderModules.R" "styles.css"
```

```r
dir("..")
```

```
## [1] "data"       "docs"      "hw"        "index.html" "index.Rmd"
## [6] "labs"       "modules"   "pdf"       "README.md"
```

Working Directory

- Copy the code to set your working directory from the History tab in RStudio (top right)
- Confirm the directory contains "day2.R" using `dir()`

Data Input

The `read.table()` function returns a `data.frame`

```r
class(mon)
```

```
## [1] "data.frame"
```

```r
str(mon)
```

```
## 'data.frame':    84 obs. of  6 variables:
##  $ name          : chr  "James Cardinal Gibbons" "The Battle Monument" "Negro Heroes of the U.S Monu
##  $ zipCode       : int  21201 21202 21202 21202 21202 21202 21202 21211 21213 21211 ...
##  $ neighborhood  : chr  "Downtown" "Downtown" "Downtown" "Downtown" ...
##  $ councilDistrict: int  11 11 11 11 11 11 11 7 14 14 ...
##  $ policeDistrict : chr  "CENTRAL" "CENTRAL" "CENTRAL" "CENTRAL" ...
##  $ Location.1     : chr  "408 CHARLES ST\nBaltimore, MD\n" "" "" "100 HOLLIDAY ST\nBaltimore, MD\n"
```

Data Input

Changing variable names in `data.frame`s works using the `names()` function, which is analagous to `colnames()` for data frames (they can be used interchangeably)

```r
names(mon)[1] = "Name"
names(mon)
```

```
## [1] "Name"          "zipCode"        "neighborhood"    "councilDistrict"
## [5] "policeDistrict" "Location.1"
```

```
names(mon)[1] = "name"
names(mon)
```

```
## [1] "name"          "zipCode"        "neighborhood"    "councilDistrict"
## [5] "policeDistrict" "Location.1"
```

Data Subsetting

Now we will introduce subsetting rows/observations of data using logical statements. Recall that the `logical` class consists of either `TRUE` or `FALSE`

```
z = c(TRUE,FALSE,TRUE,FALSE)
class(z)
```

```
## [1] "logical"
```

```
sum(z) # number of TRUEs
```

```
## [1] 2
```

Data Subsetting

And recall again that the logical class does NOT use quotes.

```
z2 = c("TRUE","FALSE","TRUE","FALSE")
class(z2)
```

```
## [1] "character"
```

```
# sum(z2)
identical(z,z2)
```

```
## [1] FALSE
```

Useful: `identical()` checks if two R objects are exactly identical/equal.

Logical Statements

Almost every R object can be evaluated and converted to the `logical` class using different logical statements (this mirrors computer science/programming syntax)

- '==': equal to
- '!=': not equal to (it is NOT '~' in R, e.g. SAS)
- '>': greater than
- '<': less than
- '>=': greater than or equal to
- '<=': less than or equal to

Logical Statements

```
x = 1:6
x > 4
```

```
## [1] FALSE FALSE FALSE FALSE  TRUE  TRUE
```

```
x == 3
```

```
## [1] FALSE FALSE  TRUE FALSE FALSE FALSE
```

Logical Statements

These logical statements can be then used to subset your data.

```
Index = (mon$zipCode == 21202)
sum(Index)
```

```
## [1] 16
```

```
table(Index)
```

```
## Index
## FALSE  TRUE
##    68    16
```

```
mon2 = mon[Index,]
```

Logical Statements

```
dim(mon2)
```

```
## [1] 16  6
```

```
head(mon2)
```

```
##                                      name zipCode neighborhood
## 2                     The Battle Monument   21202     Downtown
## 3        Negro Heroes of the U.S Monument   21202     Downtown
## 4                      Star Bangled Banner   21202     Downtown
## 5         Flame at the Holocaust Monument   21202     Downtown
## 6                           Calvert Statue   21202     Downtown
## 7 War Memorial Building/Aquatic Wa Horses   21202     Downtown
##   councilDistrict policeDistrict                     Location.1
## 2              11        CENTRAL
## 3              11        CENTRAL
## 4              11        CENTRAL 100 HOLLIDAY ST\nBaltimore, MD\n
## 5              11        CENTRAL    50 MARKET PL\nBaltimore, MD\n
## 6              11        CENTRAL  100 CALVERT ST\nBaltimore, MD\n
## 7              11        CENTRAL     101 GAY ST\nBaltimore, MD\n
```

Which

`which()`: "Give the `TRUE` indices of a logical object, allowing for array indices."

```
mon$Location.1 != ""
```

```
##  [1]  TRUE FALSE FALSE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE FALSE  TRUE
## [12] FALSE FALSE  TRUE  TRUE FALSE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE
## [23]  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE FALSE  TRUE  TRUE  TRUE
## [34]  TRUE  TRUE  TRUE  TRUE  TRUE FALSE FALSE  TRUE  TRUE  TRUE  TRUE
## [45]  TRUE  TRUE  TRUE FALSE FALSE  TRUE FALSE FALSE FALSE  TRUE  TRUE
## [56] FALSE  TRUE  TRUE  TRUE  TRUE  TRUE FALSE FALSE FALSE FALSE FALSE
## [67] FALSE  TRUE  TRUE  TRUE  TRUE  TRUE  TRUE FALSE FALSE  TRUE FALSE
## [78]  TRUE  TRUE  TRUE  TRUE FALSE FALSE  TRUE
```

```
which(mon$Location.1 != "")
```

```
##  [1]  1  4  5  6  7  8  9 11 14 15 17 18 19 20 21 22 23 24 25 26 27 28 29
## [24] 31 32 33 34 35 36 37 38 41 42 43 44 45 46 47 50 54 55 57 58 59 60 61
## [47] 68 69 70 71 72 73 76 78 79 80 81 84
```

Missing Data

- In R, missing data is represented by the symbol `NA` (note that it is NOT a character, and therefore not in quotes, just like the `logical` class)

- `is.na()` is a logical test for which variables are missing

- Many summarization functions do not the calculation you expect (e.g. they return `NA`) if there is ANY missing data, and these ofen have an argument `na.rm=FALSE`. Changing this to `na.rm=TRUE` will ignore the missing values in the calculation (i.e. `mean()`, `median()`, `max()`, `sum()`)

Here is a good link with more information: http://www.statmethods.net/input/missingdata.html