Data Cleaning

Module 9

Andrew Jaffe

June 17, 2015

Data

- We will be using multiple data sets in this lecture:
    - Salary, Monument, Circulator, and Restaurant from OpenBaltimore: https://data.baltimorecity.gov/browse?limitTo=datasets
    - Gap Minder - very interesting way of viewing longitudinal data
        * Data is here - http://www.gapminder.org/data/
    - http://spreadsheets.google.com/pub?key=rMsQHawTObBb6_U2ESjKXYw&output=xls

Data Cleaning

In general, data cleaning is a process of investigating your data for inaccuracies, or recoding it in a way that makes it more manageable.

MOST IMPORTANT RULE - LOOK AT YOUR DATA!

Again - `table`, `summarize`, `is.na`, `any`, `all` are useful.

Data Cleaning

```r
table(c(0, 1, 2, 3, NA, 3, 3, 2,2, 3),
        useNA="ifany")
```

```
##
##    0    1    2    3 <NA>
##    1    1    3    4    1
```

```r
table(c(0, 1, 2, 3, 2, 3, 3, 2,2, 3),
        useNA="always")
```

```
##
##    0    1    2    3 <NA>
##    1    1    4    4    0
```

```r
tab <- table(c(0, 1, 2, 3, 2, 3, 3, 2,2, 3),
             c(0, 1, 2, 3, 2, 3, 3, 4, 4, 3),
              useNA="always")
margin.table(tab, 2)
```

```
##
##    0    1    2    3    4 <NA>
##    1    1    2    4    2    0
```

```r
prop.table(tab)
```

```
##
##          0    1    2    3    4 <NA>
##    0    0.1 0.0 0.0 0.0 0.0  0.0
##    1    0.0 0.1 0.0 0.0 0.0  0.0
##    2    0.0 0.0 0.2 0.0 0.2  0.0
##    3    0.0 0.0 0.0 0.4 0.0  0.0
##  <NA>  0.0 0.0 0.0 0.0 0.0  0.0
```

```r
prop.table(tab,1)
```

```
##
##          0    1    2    3    4 <NA>
##    0    1.0 0.0 0.0 0.0 0.0  0.0
##    1    0.0 1.0 0.0 0.0 0.0  0.0
##    2    0.0 0.0 0.5 0.0 0.5  0.0
##    3    0.0 0.0 0.0 1.0 0.0  0.0
##  <NA>
```

Download Salary FY2014 Data

https://data.baltimorecity.gov/City-Government/Baltimore-City-Employee-Salaries-FY2014/2j28-xzd7

Download as a CSV and then read it into R as the variable `Sal`

```r
Sal = read.csv("../data/Baltimore_City_Employee_Salaries_FY2014.csv",
               as.is=TRUE)
```

Data Cleaning

- `any()` - checks if there are any TRUES
- `all()` - checks if ALL are true

```r
Sal[1:4,]
```

```
##                  Name                        JobTitle AgencyID
## 1    Aaron,Keontae E               AIDE BLUE CHIP   W02200
## 2   Aaron,Patricia G Facilities/Office Services II   A03031
## 3      Aaron,Petra L   ASSISTANT STATE'S ATTORNEY   A29005
## 4 Abaineh,Yohannes T              EPIDEMIOLOGIST   A65026
##                  Agency   HireDate AnnualSalary  GrossPay
## 1         Youth Summer  06/10/2013   $11310.00   $873.63
## 2      OED-Employment Dev  10/24/1979   $53428.00 $52868.38
## 3 States Attorneys Office  09/25/2006   $68300.00 $67439.19
## 4  HLTH-Health Department  07/23/2009   $62000.00 $58654.74
```

```
any(is.na(Sal$Name))
```

```
## [1] FALSE
```

Example of Cleaning:

For example, let's say gender was coded as Male, M, m, Female, F, f. Using Excel to find all of these would be a matter of filtering and changing all by hand or using if statements.

In R, you can simply do something like:

```
data$gender[data$gender %in%
    c("Male", "M", "m")] <- "Male"
```

Sometimes though, it's not so simple. That's where functions that find patterns come in very useful.

```
table(gender)
```

```
## gender
##      F FeMAle FEMALE     Fm      M     Ma   mAle   Male   MaLe   MALE
##     75     82     74     89     89     79     87     89     88     95
##    Man  Woman
##     73     80
```

Find/Replace and Regular Expressions

- R can do much more than find exact matches for a whole string
- Like Perl and other languages, it can use regular expressions.
- What are regular expressions?
- Ways to search for specific strings
- Can be very complicated or simple
- Highly Useful

'Find' functions

grep: `grep`, `grepl`, `regexpr` and `gregexpr` search for matches to argument pattern within each element of a character vector: they differ in the format of and amount of detail in the results.

`grep(pattern, x, fixed=FALSE)`, where:

- pattern = character string containing a regular expression to be matched in the given character vector.

- x = a character vector where matches are sought, or an object which can be coerced by as.character to a character vector.

- If fixed=TRUE, it will do exact matching for the phrase anywhere in the vector (regular find)

'Find' functions

```r
grep("Rawlings",Sal$Name)
```

```
## [1] 13832 13833 13834 13835
```

These are the indices/elements where the pattern match occurs

grep() returns something similar to which() on a logical statement

'Find' functions

```r
grep("Rawlings",Sal$Name)
```

```
## [1] 13832 13833 13834 13835
```

```r
grep("Rawlings",Sal$Name,value=TRUE)
```

```
## [1] "Rawlings,Kellye A"          "Rawlings,MarqWell D"
## [3] "Rawlings,Paula M"           "Rawlings-Blake,Stephanie C"
```

```r
Sal[grep("Rawlings",Sal$Name),]
```

```
##                              Name             JobTitle AgencyID
## 13832           Rawlings,Kellye A EMERGENCY DISPATCHER   A40302
## 13833          Rawlings,MarqWell D      AIDE BLUE CHIP   W02384
## 13834             Rawlings,Paula M       COMMUNITY AIDE   A04015
## 13835 Rawlings-Blake,Stephanie C                MAYOR   A01001
##                   Agency   HireDate AnnualSalary   GrossPay
## 13832 M-R Info Technology 01/06/2003    $47980.00  $68426.73
## 13833         Youth Summer 06/15/2012    $11310.00    $507.50
## 13834        R&P-Recreation 12/10/2007    $19802.00   $8195.79
## 13835         Mayors Office 12/07/1995   $163365.00 $161219.24
```

grep() Options

```r
head(grep("Tajhgh",Sal$Name, value=TRUE))
```

```
## [1] "Reynold,Tajhgh J"
```

```r
grep("Jaffe",Sal$Name)
```

```
## [1] 8603
```

```r
length(grep("Jaffe",Sal$Name))
```

```
## [1] 1
```

A bit on Regular Expressions

- http://www.regular-expressions.info/reference.html
- They can use to match a large number of strings in one statement
- . matches any single character
- * means repeat as many (even if 0) more times the last character
- ? makes the last thing optional

Using Regular Expressions

- Look for any name that starts with:
    - Payne at the beginning,
    - Leonard and then an S
    - Spence then a capital C

```r
grep("Payne.*", x=Sal$Name, value=TRUE)
```

```
##  [1] "Payne El,Jackie"        "Payne Johnson,Nickole A"
##  [3] "Payne,Chanel"           "Payne,Connie T"
##  [5] "Payne,Denise I"         "Payne,Dominic R"
##  [7] "Payne,James R"          "Payne,Jasman T"
##  [9] "Payne,Joey D"           "Payne,Jordan A"
## [11] "Payne,Karen V"          "Payne,Karen V"
## [13] "Payne,Leonard S"        "Payne,Mary A"
## [15] "Payne,Micah W"          "Payne,Michael C"
## [17] "Payne,Michael N"        "Payne,Morag"
## [19] "Payne,Nora M"           "Payne,Shelley F"
```

```r
grep("Leonard.?S", x=Sal$Name, value=TRUE)
```

```
## [1] "Payne,Leonard S"      "Szumlanski,Leonard S"
```

```r
grep("Spence.*C.*", x=Sal$Name, value=TRUE)
```

```
## [1] "Greene,Spencer C"    "Spencer,Charles A"    "Spencer,Christian O"
## [4] "Spencer,Clarence W"  "Spencer,Michael C"
```

Replace

Let's say we wanted to sort the data set by Annual Salary:

```r
class(Sal$AnnualSalary)
```

```
## [1] "character"
```

```r
sort(c("1", "2", "10")) #  not sort correctly (order simply ranks the data)
```

```
## [1] "1"  "10" "2"
```

```r
order(c("1", "2", "10"))
```

```
## [1] 1 3 2
```

Replace

So we must change the annual pay into a numeric:

```r
head(as.numeric(Sal$AnnualSalary), 4)
```

```
## [1] NA NA NA NA
```

R didn't like the $ so it thought turned them all to `NA`.

`sub()` and `gsub()` can do the replacing part.

Replacing and subbing

Now we can replace the $ with nothing (used fixed=TRUE because $ means something in regular expressions):

```r
Sal$AnnualSalary <- as.numeric(gsub(pattern="$", replacement="",
                                Sal$AnnualSalary, fixed=TRUE))
Sal <- Sal[order(Sal$AnnualSalary,decreasing=TRUE), ] # use negative to sort descending
Sal[1:5, c("Name", "AnnualSalary", "JobTitle")]
```

```
##                    Name AnnualSalary              JobTitle
## 1222   Bernstein,Gregg L       238772     STATE'S ATTORNEY
## 3175    Charles,Ronnie E       200000   EXECUTIVE LEVEL III
## 985      Batts,Anthony W       193800   EXECUTIVE LEVEL III
## 1343       Black,Harry E       190000   EXECUTIVE LEVEL III
## 16352      Swift,Michael       187200 CONTRACT SERV SPEC II
```

Useful String Functions

Useful String functions

- `toupper()`, `tolower()` - uppercase or lowercase your data:
- `str_trim()` (in the `stringr` package) - will trim whitespace
- `nchar` - get the number of characters in a string
- `substr(x, start, stop)` - substrings from position start to position stop
- `strsplit(x, split)` - splits strings up - returns list!
- `paste()` - paste strings together - look at `?paste`

Paste

Paste can be very useful for joining vectors together:

```r
paste("Visit", 1:5, sep="_")
```

```
## [1] "Visit_1" "Visit_2" "Visit_3" "Visit_4" "Visit_5"
```

```r
paste("Visit", 1:5, sep="_", collapse=" ")
```

```
## [1] "Visit_1 Visit_2 Visit_3 Visit_4 Visit_5"
```

```r
paste("To", "is going be the ", "we go to the store!", sep="day ")
```

```
## [1] "Today is going be the day we go to the store!"
```

```r
# and paste0 can be even simpler see ?paste0
paste0("Visit",1:5)
```

```
## [1] "Visit1" "Visit2" "Visit3" "Visit4" "Visit5"
```

```r
paste(1:5, letters[1:5], sep="_")
```

```
## [1] "1_a" "2_b" "3_c" "4_d" "5_e"
```

```r
paste(6:10, 11:15, 2000:2005, sep="/")
```

```
## [1] "6/11/2000"  "7/12/2001"  "8/13/2002"  "9/14/2003"  "10/15/2004"
## [6] "6/11/2005"
```

```r
paste(paste("x",1:5,sep=""),collapse="+")
```

```
## [1] "x1+x2+x3+x4+x5"
```

Strsplit

```r
x <- c("I really", "like writing", "R code")
y <- strsplit(x, split=" ")
y[[2]]
```

```
## [1] "like"    "writing"
```

```r
sapply(y, "[", 1) # on the fly
```

```
## [1] "I"    "like" "R"
```

```r
sapply(y, "[", 2) # on the fly
```

```
## [1] "really"  "writing" "code"
```

Data Merging/Append

- Merging - joining data sets together - usually on key variables, usually "id"
- `merge()` is the most common way to do this with data sets
- `rbind`/`cbind` - row/column bind, respectively
    - `rbind` is the equivalent of "appending" in Stata or "setting" in SAS
    - `cbind` allows you to add columns in addition to the previous ways
- `reshape2` package also has a lot of information about different ways to reshape data (wide to long, etc) - but has a different (and sometimes more intuitive syntax)
- `t()` is a function that will transpose the data

Merging

```r
base <- data.frame(id=1:10, Age= seq(55,60, length=10))
base[1:2,]
```

```
##   id      Age
## 1  1 55.00000
## 2  2 55.55556
```

```r
visits <- data.frame(id=rep(1:8, 3), visit= rep(1:3, 8),
                     Outcome= seq(10,50, length=24))
visits[1:2,]
```

```
##   id visit  Outcome
## 1  1     1 10.00000
## 2  2     2 11.73913
```

```r
merged.data <- merge(base, visits, by="id")
merged.data[1:5,]
```

```
##   id      Age visit  Outcome
## 1  1 55.00000     1 10.00000
## 2  1 55.00000     3 23.91304
## 3  1 55.00000     2 37.82609
## 4  2 55.55556     2 11.73913
## 5  2 55.55556     1 25.65217
```

```
dim(merged.data)
```

```
## [1] 24  4
```

```
all.data <- merge(base, visits, by="id", all=TRUE)
tail(all.data)
```

```
##    id      Age visit  Outcome
## 21  7 58.33333     2 48.26087
## 22  8 58.88889     2 22.17391
## 23  8 58.88889     1 36.08696
## 24  8 58.88889     3 50.00000
## 25  9 59.44444    NA       NA
## 26 10 60.00000    NA       NA
```

```
dim(all.data)
```

```
## [1] 26  4
```

Aside: Dates

You can convert date-like strings in the `Date` class (http://www.statmethods.net/input/dates.html for more info)

```
circ = read.csv("../data/Charm_City_Circulator_Ridership.csv",as.is=TRUE)
head(sort(circ$date))
```

```
## [1] "01/01/2011" "01/01/2012" "01/01/2013" "01/02/2011" "01/02/2012"
## [6] "01/02/2013"
```

```
circ$date <- as.Date(circ$date, "%m/%d/%Y") # creating a date for sorting
head(circ$date)
```

```
## [1] "2010-01-11" "2010-01-12" "2010-01-13" "2010-01-14" "2010-01-15"
## [6] "2010-01-16"
```

```
head(sort(circ$date))
```

```
## [1] "2010-01-11" "2010-01-12" "2010-01-13" "2010-01-14" "2010-01-15"
## [6] "2010-01-16"
```

Data Reshaping

Disclaimer: the `reshape` command in R is not remarkably intuitive.

- Wide - multiple measurements are variables / columns so that the data gets wider with more measurements
- Long - multiple measurements are rows so data gets longer with more measurements
- One example would be many ids with multiple visits

Example of Long/Wide

```
head(wide)
```

```
##   id visit1 visit2 visit3
## 1  1   Good   Good    Bad
```

```
head(long)
```

```
##   id visit Outcome
## 1  1     1    Good
## 2  1     2    Good
## 3  1     3     Bad
```

Data Reshaping

- Good resource: http://www.ats.ucla.edu/stat/r/faq/reshape.htm

```
head(Indometh) # this is long
```

```
## Grouped Data: conc ~ time | Subject
##   Subject time conc
## 1       1 0.25 1.50
## 2       1 0.50 0.94
## 3       1 0.75 0.78
## 4       1 1.00 0.48
## 5       1 1.25 0.37
## 6       1 2.00 0.19
```

Data Reshaping

```
wide <- reshape(Indometh, v.names = "conc", idvar = "Subject",
                timevar = "time", direction = "wide")
head(wide)
```

```
##    Subject conc.0.25 conc.0.5 conc.0.75 conc.1 conc.1.25 conc.2 conc.3
## 1        1      1.50     0.94      0.78   0.48      0.37   0.19   0.12
## 12       2      2.03     1.63      0.71   0.70      0.64   0.36   0.32
## 23       3      2.72     1.49      1.16   0.80      0.80   0.39   0.22
## 34       4      1.85     1.39      1.02   0.89      0.59   0.40   0.16
## 45       5      2.05     1.04      0.81   0.39      0.30   0.23   0.13
## 56       6      2.31     1.44      1.03   0.84      0.64   0.42   0.24
##    conc.4 conc.5 conc.6 conc.8
## 1    0.11   0.08   0.07   0.05
## 12   0.20   0.25   0.12   0.08
## 23   0.12   0.11   0.08   0.08
## 34   0.11   0.10   0.07   0.07
## 45   0.11   0.08   0.10   0.06
## 56   0.17   0.13   0.10   0.09
```

Data Reshaping

```r
dim(Indometh)
```

```
## [1] 66  3
```

```r
wide
```

```
##    Subject conc.0.25 conc.0.5 conc.0.75 conc.1 conc.1.25 conc.2 conc.3
## 1        1      1.50     0.94      0.78   0.48      0.37   0.19   0.12
## 12       2      2.03     1.63      0.71   0.70      0.64   0.36   0.32
## 23       3      2.72     1.49      1.16   0.80      0.80   0.39   0.22
## 34       4      1.85     1.39      1.02   0.89      0.59   0.40   0.16
## 45       5      2.05     1.04      0.81   0.39      0.30   0.23   0.13
## 56       6      2.31     1.44      1.03   0.84      0.64   0.42   0.24
##    conc.4 conc.5 conc.6 conc.8
## 1    0.11   0.08   0.07   0.05
## 12   0.20   0.25   0.12   0.08
## 23   0.12   0.11   0.08   0.08
## 34   0.11   0.10   0.07   0.07
## 45   0.11   0.08   0.10   0.06
## 56   0.17   0.13   0.10   0.09
```

Data Reshaping

- If you've reshaped a data set - to get it back, just reshape it again

```r
reshape(wide, direction = "long")[1:10,]
```

```
##         Subject time conc
## 1.0.25        1 0.25 1.50
## 2.0.25        2 0.25 2.03
## 3.0.25        3 0.25 2.72
## 4.0.25        4 0.25 1.85
## 5.0.25        5 0.25 2.05
## 6.0.25        6 0.25 2.31
## 1.0.5         1 0.50 0.94
## 2.0.5         2 0.50 1.63
## 3.0.5         3 0.50 1.49
## 4.0.5         4 0.50 1.39
```

Note the row name change

Data Reshaping - A Better Example

```r
TB <- read.xlsx(file="../data/indicator_estimatedincidencealltbper100000.xlsx",
                sheetName="Data")
head(TB, 1)
```

```
##    TB.incidence..all.forms..per.100.000.population.per.year. X1990 X1991
## 1                                                 Afghanistan   168   168
##   X1992 X1993 X1994 X1995 X1996 X1997 X1998 X1999 X2000 X2001 X2002 X2003
## 1   168   168   168   168   168   168   168   168   168   168   168   168
##   X2004 X2005 X2006 X2007 NA.
## 1   168   168   168   168  NA
```

```r
TB$NA. <- NULL
head(TB, 1)
```

```
##    TB.incidence..all.forms..per.100.000.population.per.year. X1990 X1991
## 1                                                 Afghanistan   168   168
##   X1992 X1993 X1994 X1995 X1996 X1997 X1998 X1999 X2000 X2001 X2002 X2003
## 1   168   168   168   168   168   168   168   168   168   168   168   168
##   X2004 X2005 X2006 X2007
## 1   168   168   168   168
```

Data Reshaping - A Better Example

```r
colnames(TB) <- c("Country", paste("Year",
                          1990:2007, sep="."))
head(TB,1)
```

```
##       Country Year.1990 Year.1991 Year.1992 Year.1993 Year.1994 Year.1995
## 1 Afghanistan       168       168       168       168       168       168
##   Year.1996 Year.1997 Year.1998 Year.1999 Year.2000 Year.2001 Year.2002
## 1       168       168       168       168       168       168       168
##   Year.2003 Year.2004 Year.2005 Year.2006 Year.2007
## 1       168       168       168       168       168
```

Data Reshaping - More is better!

```r
TB.long <- reshape(TB, idvar="Country",
          v.names="Cases", times=1990:2007,
                  direction="long", timevar="Year",
                  varying = paste("Year", 1990:2007, sep="."))
head(TB.long, 4)
```

```
##                         Country Year Cases
## Afghanistan.1990    Afghanistan 1990   168
## Albania.1990            Albania 1990    25
## Algeria.1990            Algeria 1990    38
## American Samoa.1990 American Samoa 1990    21
```

```r
rownames(TB.long) <- NULL
head(TB.long, 4)
```

```
##           Country Year Cases
## 1     Afghanistan 1990   168
## 2          Albania 1990    25
## 3          Algeria 1990    38
## 4 American Samoa 1990    21
```